

Sentiment Analysis on Amazon Fine Food Reviews

Mt. SAC CISB 62 Final Project - Fall 2023

Russell Luna

https://github.com/rmoon64/CISB62_Final/
(https://github.com/rmoon64/CISB62_Final/)

Introduction / Summary:

The 'Amazon Fine Food Reviews' from Kaggle is a complete dataset that has about 560,000 reviews, which we trim down to 50,000 reviews. The goal of the project is to recommend new products to someone who does reviews often (Chris), using RNN to perform sentimental analysis using only Amazon reviews as the model.

This project utilized various deep learning techniques. The main tools used were RNN, TensorFlow V1, and hyperparameter tuning. First, I utilized a trimmed-down version of the data set with 50,000 reviews. Then, I used Matplotlib to plot the visualizations of important data, for example, the top 10 products with the most reviewers and the percentage of positive and negative reviews made by Chris. I then prepared the RNN with the LSTM model. This model was then fitted using six epochs and a batch size of twenty. This proved to be very beneficial, as the accuracy was extremely high at 0.9953 with a quite low loss of 0.0160. Next, I cleaned the data and defined the parameters and placeholders. Lastly, I got the scores and recommendations for Chris.

<https://www.kaggle.com/datasets/snap/amazon-fine-food-reviews>
(<https://www.kaggle.com/datasets/snap/amazon-fine-food-reviews>)

```
In [1]: ▶ import tensorflow as tf
import seaborn as sns
import re
import pandas as pd
import numpy as np
import keras
import datetime

from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.preprocessing import sequence
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, Embedding
from tensorflow.keras.callbacks import ModelCheckpoint
from string import punctuation
from sklearn.model_selection import train_test_split
from nltk.corpus import stopwords
from matplotlib import pyplot as plt
from keras.preprocessing.text import Tokenizer
```

WARNING:tensorflow:From C:\Anaconda\Lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

```
In [2]: ▶ import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: ▶ #Import dataset
df = pd.read_csv('Reviews.csv')
```

```
In [4]: ▶ df.describe()
```

Out[4]:

	Id	HelpfulnessNumerator	HelpfulnessDenominator	Score	
count	50000.000000	50000.000000	50000.000000	50000.000000	5.00000
mean	25000.500000	1.603360	2.060260	4.145840	1.29519
std	14433.901067	5.620771	6.216044	1.325596	4.73462
min	1.000000	0.000000	0.000000	1.000000	9.61718
25%	12500.750000	0.000000	0.000000	4.000000	1.26964
50%	25000.500000	0.000000	1.000000	5.000000	1.30913
75%	37500.250000	2.000000	2.000000	5.000000	1.33107
max	50000.000000	398.000000	401.000000	5.000000	1.35121

In [5]: `df.head(3)`

Out[5]:

		Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
--	--	----	-----------	--------	-------------	----------------------	------------------------

0	1	B001E4KFG0	A3SGXH7AUHU8GW		delmartian	1	
1	2	B00813GRG4	A1D87F6ZCVE5NK		dll pa	0	
2	3	B000LQOCH0	ABXLMWJIXXAIN		Natalia Corres "Natalia Corres"	1	

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    50000 non-null  int64
1   ProductId            50000 non-null  object
2   UserId               50000 non-null  object
3   ProfileName          49995 non-null  object
4   HelpfulnessNumerator  50000 non-null  int64
5   HelpfulnessDenominator 50000 non-null  int64
6   Score                50000 non-null  int64
7   Time                 50000 non-null  int64
8   Summary              49998 non-null  object
9   Text                 50000 non-null  object
dtypes: int64(5), object(5)
memory usage: 3.8+ MB
```

In [7]: `print(df.shape)`

(50000, 10)

```
In [8]: df.head()
```

```
Out[8]:
```

		Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
--	--	----	-----------	--------	-------------	----------------------	------------------------

0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian		1	
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa		0	
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres	"Natalia Corres"	1	
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl		3	
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham	"M. Wassir"	0	



```
In [9]: df["Text"][0]
```

```
Out[9]: 'I have bought several of the Vitality canned dog food products and have found them all to be of good quality. The product looks more like a stew than a processed meat and it smells better. My Labrador is finicky and s he appreciates this product better than most.'
```

```
In [10]: #Grabbing the columns that we need
df_reviews = df[['ProductId', 'ProfileName', 'Score', 'Text', 'HelpfulnessNume
```

```
In [11]: ▶ #Loading the top 10 reviewers with the most reviews.  
df.pivot_table(columns=['ProfileName'], aggfunc='size').sort_values(ascending=True)
```

```
Out[11]: ProfileName  
Gary Peterson          44  
C. F. Hill "CFH"       39  
O. Brown "Ms. O. Khannah-Brown" 35  
c2                     35  
Gunner                 31  
Chris                  31  
Rebecca of Amazon "The Rebecca Review" 29  
Laura                  24  
Amanda                 24  
Mike                   24  
dtype: int64
```

```
In [12]: ▶ popular = df[['ProductId', 'ProfileName']].groupby('ProductId').count().sort_values(ascending=True)
```



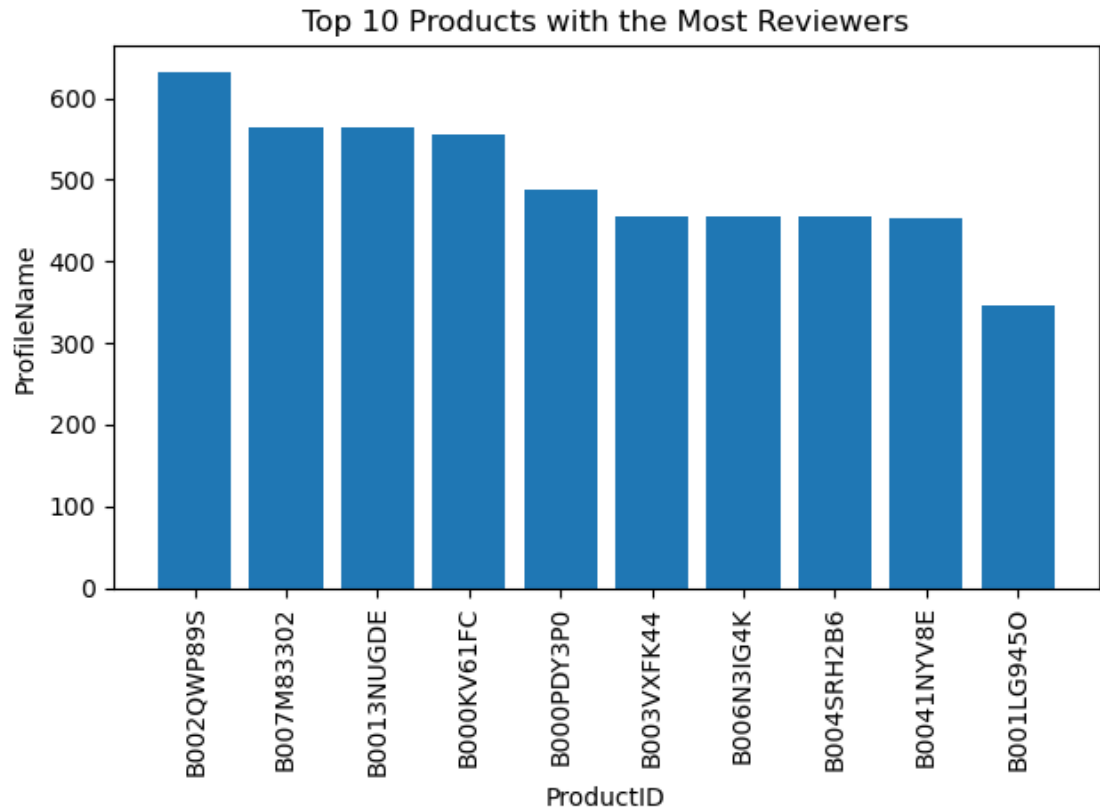
```
In [13]: ▶ df.columns
```

```
Out[13]: Index(['Id', 'ProductId', 'UserId', 'ProfileName', 'HelpfulnessNumerator',  
               'HelpfulnessDenominator', 'Score', 'Time', 'Summary', 'Text'],  
              dtype='object')
```

Visuals: matplotlib and seaborn

```
In [14]: ▶ top_10_popular = popular.head(10) # Select the top 10 products with the most reviews
```

```
In [15]: fig, axs = plt.subplots(sharey=True, tight_layout=True)
axs.bar(top_10_popular['ProductID'], top_10_popular['ProfileName'])
axs.set_xlabel('ProductID')
axs.set_ylabel('ProfileName')
axs.set_title('Top 10 Products with the Most Reviewers')
plt.xticks(rotation=90) # Rotate x-axis labels for better visibility
plt.show()
```



```
In [16]: ▶ df_reviews[df_reviews['ProfileName']=='Chris']
```


Out[16]:

	ProductId	ProfileName	Score	Text	Help
151	B00374XSVY	Chris	5	Works with chicken fish beef or pork. Fast eas...	
637	B000G6RYNE	Chris	4	This kettle chips taste "Good , Crispy & Crunc...	
729	B008BEGP9W	Chris	4	I really like the pineapple shortcakes sold he...	
1006	B002XG21MO	Chris	5	These are just like the animal crackers we eat...	
3142	B000FDKQCY	Chris	5	This product is as good as any that I have eve...	
5301	B003OJLCXI	Chris	4	I really like this tea, but my husband smelled...	
5826	B000633O2C	Chris	5	I want to give my cats as much variety as poss...	
7424	B005LMLXN0	Chris	5	As a tea-junkie, Runa's traditional guayusa of...	
7653	B000TRFGGM	Chris	4	I opened the carton and the plant looked like ...	
7812	B00474H936	Chris	4	I have tried almost every kind of "fake meat" ...	
9712	B001LXYA5Q	Chris	4	This stuff works, but be warned the effect doe...	
14414	B00063KO34	Chris	5	My dogs love these chews. I have 5 dogs - all ...	
15075	B00113L7KC	Chris	5	I love this tea. I don't profess to be a tea ...	
15150	B0013E21V8	Chris	5	Our family loves this syrup! I love that it i...	
16240	B007TJGZ54	Chris	5	Love the Green Mountain Coffee; tastes great i...	
16546	B003P7WLF2	Chris	1	I purchased these in lieu of the V8 Fusions si...	
18674	B00008JOL0	Chris	5	Zuke's is the best and my dog loves to eat the...	
21217	B002QWP89S	Chris	5	Molly and Cocoa, our cocker spaniels, said to ...	
21607	B002QWP89S	Chris	5	I am happy with the service of this company. ...	
26212	B0058AMY74	Chris	4	This kettle chips taste "Good , Crispy & Crunc...	
26717	B001IZG6P4	Chris	5	These are the best peppercorns for the money I...	
30041	B000DZFMEQ	Chris	5	My wife is gluten and dairy intolerant and we ...	
32735	B0083QJU72	Chris	3	Having grown up in sugar maple country and bou...	
36181	B000ILKXP8	Chris	3	I've seen some reviewers commenting on the sug...	
38414	B002QK0RJG	Chris	5	<a href="http://www.amazon.com/gp/product/B004...	
41755	B0000D8DI0	Chris	4	This is excellent chocolate. Hachez is a qual...	
43455	B001EQ4P2I	Chris	5	I'm a big fan of Blue Diamond brand Almonds. I...	
43457	B001EQ4P2I	Chris	5	I'm a big fan of the Blue Diamond brand almond...	
43524	B001EQ4P2I	Chris	3	To start, I'm a big fan of Blue Diamond brand ...	
46281	B001EO5KSS	Chris	4	These are great! Our whole family likes them....	
49031	B000X1F9XY	Chris	1	2 of my 3 dogs destroyed these in less than 30...	

Plotting Chris' percentage of positive views vs negative reviews


```
In [17]: df_chris = df_reviews[df_reviews['ProfileName']=='Chris']
```

```
In [18]: fig, axs = plt.subplots(sharey=True, tight_layout=True)
axs.hist(df_chris['Score'], bins=5)
plt.xlim([1, 5])
plt.xticks(np.arange(1, 6, 1))
axs.set_title('Percentage of Positive and Negative Reviews for Chris')
axs.set_xlabel('Score')
axs.set_ylabel('Frequency')
plt.show()
```



Plotting the total reviews

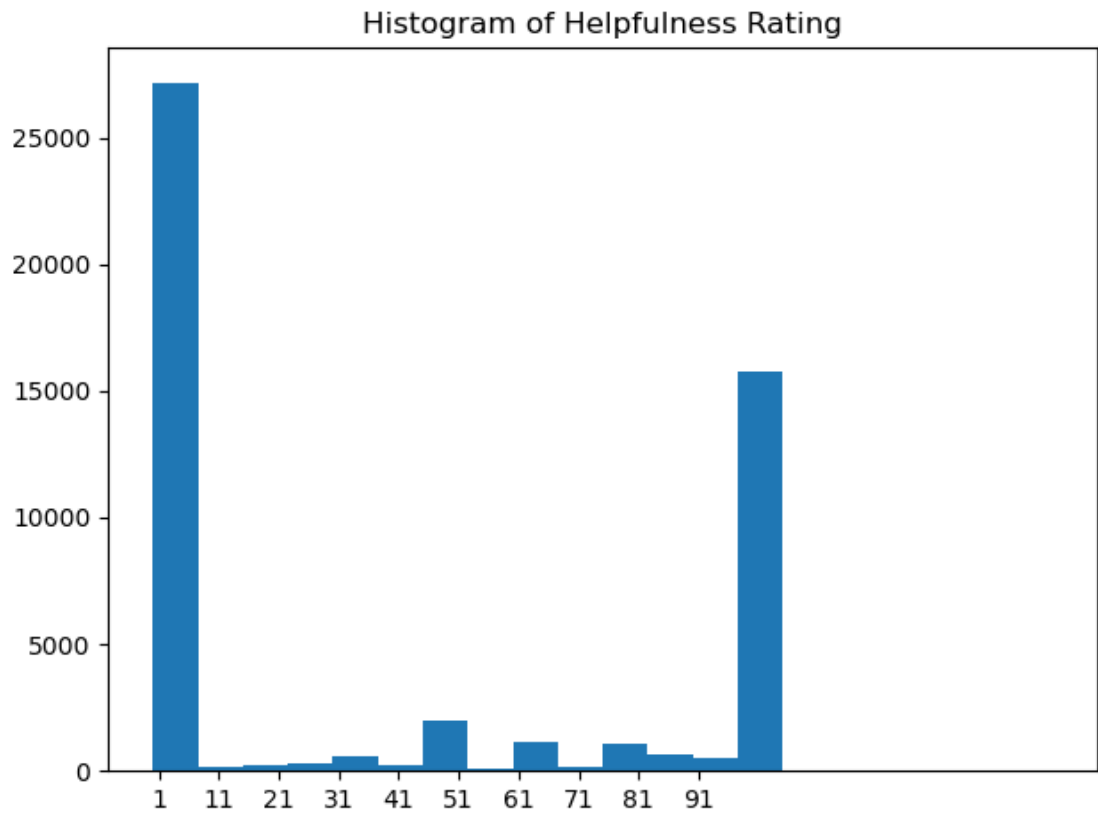
```
In [19]: ▶ fig, axs = plt.subplots(sharey=True, tight_layout=True)
          axs.hist(df_reviews['Score'], bins=5)
          plt.xlim([1, 5])
          plt.xticks(np.arange(1, 6, 1))
          plt.title('Total Reviews')
          plt.show()
```



Histogram of helpfulness rating

```
In [20]: ▶ df_reviews['Percent'] = (df_reviews['HelpfulnessNumerator'] / df_reviews['
          df_reviews['Percent'].fillna(0,inplace=True)
```


```
In [21]: fig, axs = plt.subplots(sharey=True, tight_layout=True)
axs.hist(df_reviews['Percent'], bins=20)
plt.xticks(np.arange(1, 100, 10))
plt.title('Histogram of Helpfulness Rating')
plt.show()
```



```
In [22]: df_reviews['Polarity_Rating'] = df_reviews['Score'].apply(lambda x: 'Posit
```

```
In [23]: sns.countplot(data = df_reviews, x= 'Polarity_Rating')
```

```
Out[23]: <Axes: xlabel='Polarity_Rating', ylabel='count'>
```



```
In [24]: df_reviews[df_reviews['Polarity_Rating'] == 'Negative'].shape[0]
```

```
Out[24]: 11582
```

```
In [25]: data_Positive = df_reviews[df_reviews['Polarity_Rating'] == 'Positive'][0:]
data_Negative = df_reviews[df_reviews['Polarity_Rating'] == 'Negative']

data_Negative_over = data_Negative.sample(20000, replace=True)
df_balance_reviews = pd.concat([data_Positive, data_Negative_over], axis=0)
```

```
In [26]: import nltk
nltk.download('stopwords')
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\RL\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
Out[26]: True
```

```
In [27]: #Setting the stop words
english_stops = set(stopwords.words('english'))
```

```
In [28]: ▶ df = df_balance_reviews
x_data = df['Text']          # Reviews/Input
y_data = df['Polarity_Rating'] # Sentiment/Output

# PRE-PROCESS REVIEW
x_data = x_data.replace({'<.*?>': ''}, regex = True)      # remove html
x_data = x_data.replace({'[^A-Za-z]': ' '}, regex = True) # remove non-alphabetic
x_data = x_data.apply(lambda review: [w for w in review.split() if w not in stopwords], axis=1)
x_data = x_data.apply(lambda review: [w.lower() for w in review], axis=1) # Lowercase

# ENCODE SENTIMENT -> 0 & 1
y_data = y_data.replace('Positive', 1)
### similarly for y_data
y_data = y_data.replace('Negative', 0)
```

```
In [29]: ▶ X_train, X_test, y_train, y_test = train_test_split(x_data, y_data, test_size=0.2)
```

```
In [30]: ▶ def get_max_length(x_train):
    review_length = []
    for review in x_train:
        review_length.append(len(review))

    return int(np.ceil(np.mean(review_length)))
```

```
In [31]: ▶ max_length = get_max_length(X_train)
print(max_length)
```

48

```
In [32]: ▶ # ENCODE REVIEW
token = Tokenizer(lower=False) # False because we already did it.
token.fit_on_texts(X_train)
x_train = token.texts_to_sequences(X_train)
x_test = token.texts_to_sequences(X_test)

### this is the code to add the pad to x_train
x_train = pad_sequences(x_train, maxlen=max_length, padding='post', truncate='post')

##### add pad to x_test #####
x_test = pad_sequences(x_test, maxlen=max_length, padding='post', truncate='post')

total_words = len(token.word_index) + 1 # we need to add 1 because of 0
```

```
In [33]: ▶ # ARCHITECTURE
EMBED_DIM = 3200
LSTM_OUT = 64

#### add the model here:
model = Sequential()

#model.add(Embedding(input_dim=max_length, output_dim=LSTM_OUT, input_length=
model.add(Embedding(total_words, EMBED_DIM, input_length = max_length))
model.add(LSTM(LSTM_OUT))
model.add(Dense(units=1, activation='sigmoid'))

### compile the model using: optimizer = 'adam', Loss = 'binary_crossentropy'
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accu

#Adding a Checkpoint to save the model.
checkpoint = ModelCheckpoint('models/LSTM.h5', monitor='accuracy', save_be
```

WARNING:tensorflow:From C:\Anaconda\Lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From C:\Anaconda\Lib\site-packages\keras\src\optimizers__init__.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

```
In [34]: ▶ model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 48, 3200)	84083200
lstm (LSTM)	(None, 64)	835840
dense (Dense)	(None, 1)	65
=====		
Total params: 84919105 (323.94 MB)		
Trainable params: 84919105 (323.94 MB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [35]: ▶ # Fit the model
history = model.fit(x_train, y_train, epochs=6, batch_size=20, callbacks=
```


Epoch 1/6
WARNING:tensorflow:From C:\Anaconda\Lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From C:\Anaconda\Lib\site-packages\keras\src\engine\base_layer_utils.py:384: The name tf.executing_eagerly_outside_functions is deprecated. Please use tf.compat.v1.executing_eagerly_outside_functions instead.


1300/1300 [=====] - ETA: 0s - loss: 0.4488 - accuracy: 0.7871
Epoch 1: accuracy improved from -inf to 0.78708, saving model to models\LSTM.h5
1300/1300 [=====] - 722s 550ms/step - loss: 0.4488 - accuracy: 0.7871 - val_loss: 0.3165 - val_accuracy: 0.8706
Epoch 2/6
1300/1300 [=====] - ETA: 0s - loss: 0.1882 - accuracy: 0.9298
Epoch 2: accuracy improved from 0.78708 to 0.92981, saving model to models\LSTM.h5
1300/1300 [=====] - 707s 544ms/step - loss: 0.1882 - accuracy: 0.9298 - val_loss: 0.2889 - val_accuracy: 0.8880
Epoch 3/6
1300/1300 [=====] - ETA: 0s - loss: 0.0871 - accuracy: 0.9716
Epoch 3: accuracy improved from 0.92981 to 0.97158, saving model to models\LSTM.h5
1300/1300 [=====] - 834s 641ms/step - loss: 0.0871 - accuracy: 0.9716 - val_loss: 0.3048 - val_accuracy: 0.9006
Epoch 4/6
1300/1300 [=====] - ETA: 0s - loss: 0.0451 - accuracy: 0.9853
Epoch 4: accuracy improved from 0.97158 to 0.98531, saving model to models\LSTM.h5
1300/1300 [=====] - 902s 694ms/step - loss: 0.0451 - accuracy: 0.9853 - val_loss: 0.4040 - val_accuracy: 0.8922
Epoch 5/6
1300/1300 [=====] - ETA: 0s - loss: 0.0220 - accuracy: 0.9935
Epoch 5: accuracy improved from 0.98531 to 0.99354, saving model to models\LSTM.h5
1300/1300 [=====] - 801s 616ms/step - loss: 0.0220 - accuracy: 0.9935 - val_loss: 0.4108 - val_accuracy: 0.9106
Epoch 6/6
1300/1300 [=====] - ETA: 0s - loss: 0.0160 - accuracy: 0.9953
Epoch 6: accuracy improved from 0.99354 to 0.99527, saving model to models\LSTM.h5
1300/1300 [=====] - 775s 596ms/step - loss: 0.0160 - accuracy: 0.9953 - val_loss: 0.4495 - val_accuracy: 0.9083

```
In [36]: ▶ plt.figure(figsize=(12,6))
plt.plot(model.history.history['loss'])
plt.plot(model.history.history['val_loss'])
plt.title('Model loss')
plt.xlabel('epoch')
plt.ylabel('loss (categorical_crossentropy)')
plt.legend(['Train', 'Val'], loc='upper right')
```

Out[36]: <matplotlib.legend.Legend at 0x238c30a1450>




```
In [37]: ▶ plt.figure(figsize=(12,6))
plt.plot(model.history.history['accuracy'][:])
plt.plot(model.history.history['val_accuracy'][:])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.show()
```



```
In [38]: ▶ #I grab a random review from the dataframe
value = 200
df_reviews.iloc[value]
```

```
Out[38]: ProductId      B0028C4
4Z0
ProfileName
Lou
Score
2
Text      Even with small containers, they don't fill t
h...
HelpfulnessNumerator
0
HelpfulnessDenominator
0
Percent
0.0
Polarity_Rating      Negat
ive
Name: 200, dtype: object
```

```
In [39]: df_reviews.iloc[value]['Text']
```

```
Out[39]: "Even with small containers, they don't fill them up. These little tins  
are less than half filled and at the price charged it seems a rip-off. I  
s there some exotic ingredient as costly as gold contained in those tiny  
squares? Or how about the cereal ploy, they were filled at the factory  
but settled in transport.<br />Can manufacturers be honest in their deal  
ings?"
```

```
In [40]: review = df_reviews.iloc[value]['Text']  
  
regex = re.compile(r'^a-zA-Z\s')  
review = regex.sub('', review)  
  
words = review.split(' ')  
filtered = [w for w in words if w not in english_stops]  
filtered = ' '.join(filtered)  
filtered = [filtered.lower()]
```

```
In [41]: tokenize_words = token.texts_to_sequences(filtered)  
tokenize_words = pad_sequences(tokenize_words, maxlen=max_length, padding=
```

```
In [42]: result = model.predict(tokenize_words)  
  
if result >= .50:  
    print('Positive')  
else:  
    print('Negative')  
  
1/1 [=====] - 1s 631ms/step  
Negative
```

```
In [43]: #Making an Unique Numeric ID for all the users  
df_reviews['UserID'] = pd.factorize(df_reviews['ProfileName'])[0] + 1
```

```
In [44]: df_reviews[df_reviews['UserID'] == 151].head()
```

```
Out[44]:
```

	ProductId	ProfileName	Score	Text	HelpfulnessNumerator	HelpfulnessDenominator
151	B00374XSVY	Chris	5	Works with chicken fish beef or pork. Fast eas...	0	
637	B000G6RYNE	Chris	4	This kettle chips taste "Good , Crispy & Crunc...	0	
729	B008BEGP9W	Chris	4	I really like the pineapple shortcakes sold he...	0	
1006	B002XG21MO	Chris	5	These are just like the animal crackers we eat...	0	
3142	B000FDKQCY	Chris	5	This product is as good as any that I have eve...	0	

```
In [45]: #I had to clean the data some more, some review went back and review the s
cleaned = df_reviews.drop_duplicates(subset=['ProductId', 'UserID'], keep=

print('Removed Dulipcated, Total Reviews {}'.format(cleaned.shape[0]))
print('Total Reviews {}'.format(df_reviews.shape[0]))
```

```
Removed Dulipcated, Total Reviews 49283
Total Reviews 50000
```

In [46]: `cleaned.head()`

Out[46]:

	ProductId	ProfileName	Score	Text	HelpfulnessNumerator	HelpfulnessDenomin
--	-----------	-------------	-------	------	----------------------	--------------------

0	B001E4KFG0	delmartian	5	I have bought several of the Vitality canned d...	1	
1	B00813GRG4	dll pa	1	Product arrived labeled as Jumbo Salted Peanut...	0	
2	B000LQOCH0	Natalia Corres "Natalia Corres"	4	This is a confection that has been around a fe...	1	
3	B000UA0QIQ	Karl	2	If you are looking for the secret ingredient i...	3	
4	B006K2ZZ7K	Michael D. Bigham "M. Wassir"	5	Great taffy at a great price. There was a wid...	0	

In [47]: `user_rating_df = cleaned.pivot(index='UserID', columns='ProductId', values=`

In [48]: `user_rating_df.head(5)`

Out[48]:

	ProductId	2734888454	B00002NCJC	B00002Z754	B00005V3DC	B000084DVR	B000084E1U
--	-----------	------------	------------	------------	------------	------------	------------

UserID							
0		NaN	NaN	NaN	NaN	NaN	NaN
1		NaN	NaN	NaN	NaN	NaN	NaN
2		NaN	NaN	NaN	NaN	NaN	NaN
3		NaN	NaN	NaN	NaN	NaN	NaN
4		NaN	NaN	NaN	NaN	NaN	NaN

5 rows × 6115 columns

```
In [49]: ▶ tf.compat.v1.disable_eager_execution()
```

```
In [50]: ▶ norm_user_rating_df = user_rating_df.fillna(0) / 5.0  
trX = norm_user_rating_df.values  
trX[0:5]
```

```
Out[50]: array([[0., 0., 0., ..., 0., 0., 0.],  
                [0., 0., 0., ..., 0., 0., 0.],  
                [0., 0., 0., ..., 0., 0., 0.],  
                [0., 0., 0., ..., 0., 0., 0.],  
                [0., 0., 0., ..., 0., 0., 0.]])
```



```

In [51]: ▶ hiddenUnits = 20
visibleUnits = len(user_rating_df.columns)
vb = tf.compat.v1.placeholder("float", [visibleUnits])
hb = tf.compat.v1.placeholder("float", [hiddenUnits])
W = tf.compat.v1.placeholder("float", [visibleUnits, hiddenUnits])

# Process phase 1 of a RBM, use v0, _h0, h0
v0 = tf.compat.v1.placeholder("float", [None, visibleUnits])
_h0 = tf.nn.sigmoid(tf.matmul(v0, W) + hb)
h0 = tf.nn.relu(tf.sign(_h0 - tf.random.uniform(tf.shape(input=_h0))))

#Phase 2: Reconstruction
_v1 = tf.nn.sigmoid(tf.matmul(h0, tf.transpose(W)) + vb)
v1 = tf.nn.relu(tf.sign(_v1 - tf.random.uniform(tf.shape(input=_v1))))
h1 = tf.nn.sigmoid(tf.matmul(v1, W) + hb)

#Learning rate
alpha = 1.0

w_pos_grad = tf.matmul(tf.transpose(v0), h0)
w_neg_grad = tf.matmul(tf.transpose(v1), h1)

CD = (w_pos_grad - w_neg_grad) / tf.compat.v1.to_float(tf.shape(v0)[0])

update_w = W + alpha * CD
update_vb = vb + alpha * tf.reduce_mean(v0 - v1, 0)
update_hb = hb + alpha * tf.reduce_mean(h0 - h1, 0)

err = v0 - v1
err_sum = tf.reduce_mean(err * err)

#Current weight
cur_w = np.zeros([visibleUnits, hiddenUnits], np.float32)

cur_vb = np.zeros([visibleUnits], np.float32)

cur_hb = np.zeros([hiddenUnits], np.float32)

prv_w = np.zeros([visibleUnits, hiddenUnits], np.float32)

prv_vb = np.zeros([visibleUnits], np.float32)

prv_hb = np.zeros([hiddenUnits], np.float32)

sess = tf.compat.v1.Session()
sess.run(tf.compat.v1.global_variables_initializer())

# Using a for Loop, Run the model, with 10 epochs, batchsize = 100, and er
epochs = 10
batchsize = 100
errors = []
for i in range(epochs):
    for start, end in zip(range(0, len(trX), batchsize), range(batchsize,
        batch = trX[start:end]
        cur_w = sess.run(update_w, feed_dict={v0: batch, W: prv_w, vb: prv
        cur_vb = sess.run(update_vb, feed_dict={v0: batch, W: prv_w, vb: p

```

```

cur_nb = sess.run(update_hb, feed_dict={v0: batch, W: prv_w, vb: p
prv_w = cur_w
prv_vb = cur_vb
prv_hb = cur_hb
errors.append(sess.run(err_sum, feed_dict={v0: trX, W: cur_w, vb: cur_
print (errors[-1])

```

WARNING:tensorflow:From C:\Anaconda\Lib\site-packages\tensorflow\python\util\dispatch.py:1260: to_float (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use `tf.cast` instead.

0.0025635296

0.0013016531

0.0009061099

0.0007227106

0.00061612227

0.00055364566

0.00051163853

0.00048007743


0.00045839624

0.0004471017

```

In [52]: ▶ plt.plot(errors)
plt.ylabel('Error')
plt.xlabel('Epoch')
plt.show()

```




```
In [53]: mock_user_id = 151
```

```
In [54]: inputUser = trX[mock_user_id-1].reshape(1, -1)
inputUser[0:5]
```

```
Out[54]: array([[0., 0., 0., ..., 0., 0., 0.]])
```

```
In [55]: # Feeding in the user and reconstructing the input. use sigmoid
hh0 = tf.nn.sigmoid(tf.matmul(v0, W) + hb)
vv1 = tf.nn.sigmoid(tf.matmul(hh0, tf.transpose(W)) + vb)
feed = sess.run(hh0, feed_dict={ v0: inputUser, W: prv_w, hb: prv_hb})
rec = sess.run(vv1, feed_dict={ hh0: feed, W: prv_w, vb: prv_vb})
```

```
In [56]: scored_asin_df_mock = cleaned.drop_duplicates(subset=['ProductId'])
scored_asin_df_mock = scored_asin_df_mock.assign(RecommendationScore = rec)
scored_asin_df_mock[['ProductId', 'RecommendationScore']].sort_values(["Rec
```

```
Out[56]:
```

	ProductId	RecommendationScore
--	-----------	---------------------

4063	B001EW5YQS	0.025583
7321	B0042395CA	0.012464
45379	B000FVVZ3U	0.008445
40981	B000IKEGRK	0.005815
36090	B002DHN1UE	0.005210
17705	B0010AR1E2	0.004573
5012	B000HDKZDC	0.004495
17240	B000I6PZVK	0.004260
14515	B0016814QO	0.004161
33278	B001RQAVKK	0.004146
29077	B000NU8H0C	0.004062
2950	B003XKKEBE	0.003626
18894	B003FSPWXY	0.003173
4701	B0048IPR9Y	0.002770
17127	B008FSIHNG	0.002756
13021	B001EPQ8XS	0.002741
35800	B00290W1CY	0.002572
13806	B000FFIIYU	0.002498
19006	B003R0LM48	0.002466
7881	B000VZJJIS	0.002433

```
In [57]: food_df_mock = cleaned[cleaned['UserID'] == 151]
         food_df_mock.head()
```

Out[57]:

	ProductId	ProfileName	Score	Text	HelpfulnessNumerator	HelpfulnessDenominator
151	B00374XSVY	Chris	5	Works with chicken fish beef or pork. Fast eas...	0	0
637	B000G6RYNE	Chris	4	This kettle chips taste "Good , Crispy & Crunc...	0	0
729	B008BEGP9W	Chris	4	I really like the pineapple shortcakes sold he...	0	0
1006	B002XG21MO	Chris	5	These are just like the animal crackers we eat...	0	0
3142	B000FDKQCY	Chris	5	This product is as good as any that I have eve...	0	0

```
In [58]: merged_df_mock = scored_asin_df_mock.merge(food_df_mock, on='ProductId', how='left')
```

```
In [59]: merged_df_mock[['ProductId', 'RecommendationScore']].sort_values(["RecommendationScore"], ascending=False)
```

Out[59]:

	ProductId	RecommendationScore
649	B001EW5YQS	0.025583
1122	B0042395CA	0.012464
5653	B000FVVZ3U	0.008445
5023	B000IKEGRK	0.005815
4366	B002DHN1UE	0.005210

Above are the top 5 items that Chris will be interested in purchasing for reviewing.

Conclusion

In this project, I successfully conducted sentiment analysis on the Amazon Fine Food Reviews dataset using Jupyter Notebooks and various natural language processing techniques. Our main objective was to recommend new products to someone who does reviews often (Chris), using RNN to perform sentiment analysis using only Amazon reviews as the model. I began by performing exploratory data analysis to gain insights into the dataset's characteristics and distribution of sentiments. Preprocessing steps, such as removing stop words, were applied to clean the text data and enhance the quality of our sentiment analysis. This analysis provided valuable insights into the key factors driving positive or negative sentiment within the Amazon Fine Food Reviews dataset.

The Sentiment Analysis on Amazon Fine Food Reviews project successfully employed deep learning techniques to classify customer reviews as positive or negative. The outcome of this project contributes to a better understanding of customer sentiments and can be potentially applied in various industries for customer feedback analysis and decision-making processes. Overall, this project was a success.